

Merceded_Benz

July 13, 2023

```
[1]: # Importing the required libraries
import numpy as np
import pandas as pd
from sklearn.decomposition import PCA
```

```
[3]: # Importing the data

train = pd.read_csv('trainMB.csv')
test = pd.read_csv('testMB.csv')
```

```
[4]: train.head()
```

```
[4]:   ID      y  X0 X1  X2 X3 X4 X5 X6 X8 ... X375  X376  X377  X378  X379  \
0   0  130.81   k  v   at  a  d  u  j  o ...    0    0    1    0    0
1   6   88.53   k  t   av  e  d  y  l  o ...    1    0    0    0    0
2   7   76.26  az  w   n  c  d  x  j  x ...    0    0    0    0    0
3   9   80.62  az  t   n  f  d  x  l  e ...    0    0    0    0    0
4  13   78.02  az  v   n  f  d  h  d  n ...    0    0    0    0    0
```

```
      X380  X382  X383  X384  X385
0         0     0     0     0     0
1         0     0     0     0     0
2         0     1     0     0     0
3         0     0     0     0     0
4         0     0     0     0     0
```

[5 rows x 378 columns]

```
[5]: test.head()
```

```
[5]:   ID  X0 X1  X2 X3 X4 X5 X6 X8  X10 ... X375  X376  X377  X378  X379  X380  \
0   1  az  v   n  f  d  t  a  w    0 ...    0    0    0    1    0    0
1   2   t  b  ai  a  d  b  g  y    0 ...    0    0    1    0    0    0
2   3  az  v  as  f  d  a  j  j    0 ...    0    0    0    1    0    0
3   4  az  l   n  f  d  z  l  n    0 ...    0    0    0    1    0    0
4   5   w  s  as  c  d  y  i  m    0 ...    1    0    0    0    0    0
```

	X382	X383	X384	X385
0	0	0	0	0
1	0	0	0	0
2	0	0	0	0
3	0	0	0	0
4	0	0	0	0

[5 rows x 377 columns]

```
[6]: print('Size of training set: {} rows and {} columns'.format(*train.shape))
      print('Size of testing set: {} rows and {} columns'.format(*test.shape))
```

Size of training set: 4209 rows and 378 columns
Size of testing set: 4209 rows and 377 columns

```
[7]: # Collect the Y values into an array
      y_train = train['y'].values
```

```
[8]: y_train
```

```
[8]: array([130.81,  88.53,  76.26, ..., 109.22,  87.48, 110.85])
```

```
[9]: # Understand the data types
      cols = [c for c in train.columns if 'X' in c]
      print('Number of features: {}'.format(len(cols)))
      print('Feature types:')
      train[cols].dtypes.value_counts()
```

Number of features: 376
Feature types:

```
[9]: int64    368
      object     8
      dtype: int64
```

```
[10]: # Count the data in each of the columns
```

```
counts = [[], [], []]
for c in cols:
    typ = train[c].dtype
    uniq = len(np.unique(train[c]))
    if uniq == 1:
        counts[0].append(c)
    elif uniq == 2 and typ == np.int64:
        counts[1].append(c)
    else:
        counts[2].append(c)
```

```
print('Constant features: {} Binary features: {} Categorical features: {}\n'
      .format(*[len(c) for c in counts]))
print('Constant features:', counts[0])
print('Categorical features:', counts[2])
```

Constant features: 12 Binary features: 356 Categorical features: 8

Constant features: ['X11', 'X93', 'X107', 'X233', 'X235', 'X268', 'X289',
'X290', 'X293', 'X297', 'X330', 'X347']

Categorical features: ['X0', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X8']

```
[11]: # Splitting the data
usable_columns = list(set(train.columns) - set(['ID', 'y']))
y_train = train['y'].values
id_test = test['ID'].values
x_train = train[usable_columns]
x_test = test[usable_columns]
```

Check for null values and unique values for train & test data

```
[12]: def check_missing_values(df):
        if df.isnull().any().any():
            print('There are missing values in the dataframe')
        else:
            print('There are no missing values in the dataframe')
```

```
[13]: check_missing_values(x_train)
        check_missing_values(x_test)
```

There are no missing values in the dataframe

There are no missing values in the dataframe

Label Encoding the categorical values

```
[14]: for column in usable_columns:
        cardinality = len(np.unique(x_train[column]))
        if cardinality == 1:
            x_train.drop(column, axis=1) # Column with only one
            # value is useless so we drop it
            x_test.drop(column, axis=1)
        if cardinality > 2: # Column is categorical
            mapper = lambda x: sum([ord(digit) for digit in x])
            x_train[column] = x_train[column].apply(mapper)
            x_test[column] = x_test[column].apply(mapper)
x_train.head()
```

/usr/local/lib/python3.7/site-packages/ipykernel_launcher.py:9:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
if __name__ == '__main__':
/usr/local/lib/python3.7/site-packages/ipykernel_launcher.py:10:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
# Remove the CWD from sys.path while we load stuff.
```

```
[14]:      X42  X50  X246  X379  X155  X118  X294  X143  X101  X22  ...  X124  X233  \
0      0    0    0      0      0      1      0      0      0      0  ...      0      0
1      0    0    0      0      0      1      0      0      1      0  ...      0      0
2      0    0    1      0      0      0      0      0      1      0  ...      0      0
3      0    0    1      0      0      0      0      0      1      0  ...      0      0
4      0    0    1      0      0      0      0      0      1      0  ...      0      0

      X225  X270  X323  X277  X275  X205  X80  X67
0      0      0      0      0      1      0      0      0
1      0      0      0      0      1      1      1      0
2      0      0      0      0      0      1      1      0
3      0      0      0      0      0      1      1      0
4      0      0      0      0      0      1      1      0
```

[5 rows x 376 columns]

```
[15]: # Make sure the data is changed into numerical values

print('Feature types:')
x_train[cols].dtypes.value_counts()
```

Feature types:

```
[15]: int64      376
dtype: int64
```

Perform Dimensionality reduction

```
[16]: n_comp = 12
pca = PCA(n_components = n_comp,random_state = 420)
pca2_results_train = pca.fit_transform(x_train)
pca2_results_test = pca.transform(x_test)
```

Training using XGBoost

```
[17]: # Training using XGBoost

import xgboost as xgb
from sklearn.metrics import r2_score
from sklearn.model_selection import train_test_split

[18]: x_train,x_val,y_train,y_val = train_test_split(pca2_results_train, y_train,
↳test_size=0.2, random_state=4242)

[19]: d_train = xgb.DMatrix(x_train,label = y_train)
d_val = xgb.DMatrix(x_val,label = y_val)

# dtest = xgb.DMatrix(x_test)

d_test = xgb.DMatrix(pca2_results_test)

[20]: params = {}
params['objective'] = 'reg:linear'
params['eta'] = 0.02
params['max_depth'] = 4

def xgb_r2_score(preds, dtrain):
    labels = dtrain.get_label()
    return 'r2', r2_score(labels, preds)
watchlist = [(d_train, 'train'), (d_val, 'valid')]
clf = xgb.train(params, d_train, 1000, watchlist, early_stopping_rounds=50,
    feval=xgb_r2_score, maximize=True, verbose_eval=10)
```

[12:48:41] WARNING: /workspace/src/objective/regression_obj.cu:167: reg:linear is now deprecated in favor of reg:squarederror.

```
[0]      train-rmse:99.14835      valid-rmse:98.26297      train-r2:-58.35295
valid-r2:-67.63754
```

Multiple eval metrics have been passed: 'valid-r2' will be used for early stopping.

Will train until valid-r2 hasn't improved in 50 rounds.

```
[10]      train-rmse:81.27653      valid-rmse:80.36433      train-r2:-38.88428
valid-r2:-44.91014
```

```
[20]      train-rmse:66.71610      valid-rmse:65.77334      train-r2:-25.87403
valid-r2:-29.75260
```

```
[30]      train-rmse:54.86915      valid-rmse:53.89120      train-r2:-17.17724
valid-r2:-19.64513
```

```
[40]      train-rmse:45.24564      valid-rmse:44.22231      train-r2:-11.36018
valid-r2:-12.90160
```

```
[50]      train-rmse:37.44742      valid-rmse:36.37758      train-r2:-7.46672
valid-r2:-8.40697
```

```
[60]      train-rmse:31.15106      valid-rmse:30.01765      train-r2:-4.85891
```

valid-r2:-5.40524		
[70] train-rmse:26.08771	valid-rmse:24.90846	train-r2:-3.10907
valid-r2:-3.41038		
[80] train-rmse:22.04900	valid-rmse:20.82502	train-r2:-1.93528
valid-r2:-2.08285		
[90] train-rmse:18.84765	valid-rmse:17.60010	train-r2:-1.14480
valid-r2:-1.20198		
[100] train-rmse:16.33699	valid-rmse:15.08526	train-r2:-0.61145
valid-r2:-0.61766		
[110] train-rmse:14.39788	valid-rmse:13.15610	train-r2:-0.25161
valid-r2:-0.23037		
[120] train-rmse:12.93041	valid-rmse:11.69388	train-r2:-0.00948
valid-r2:0.02793		
[130] train-rmse:11.81665	valid-rmse:10.62117	train-r2:0.15694
valid-r2:0.19809		
[140] train-rmse:10.98570	valid-rmse:9.85576	train-r2:0.27134
valid-r2:0.30950		
[150] train-rmse:10.37823	valid-rmse:9.32776	train-r2:0.34969
valid-r2:0.38150		
[160] train-rmse:9.92529	valid-rmse:8.96124	train-r2:0.40522
valid-r2:0.42916		
[170] train-rmse:9.59273	valid-rmse:8.71470	train-r2:0.44441
valid-r2:0.46013		
[180] train-rmse:9.34136	valid-rmse:8.55182	train-r2:0.47314
valid-r2:0.48013		
[190] train-rmse:9.16018	valid-rmse:8.44863	train-r2:0.49338
valid-r2:0.49260		
[200] train-rmse:9.01539	valid-rmse:8.38342	train-r2:0.50927
valid-r2:0.50040		
[210] train-rmse:8.91053	valid-rmse:8.34375	train-r2:0.52062
valid-r2:0.50511		
[220] train-rmse:8.83439	valid-rmse:8.31880	train-r2:0.52878
valid-r2:0.50807		
[230] train-rmse:8.76692	valid-rmse:8.30589	train-r2:0.53595
valid-r2:0.50960		
[240] train-rmse:8.71890	valid-rmse:8.30186	train-r2:0.54102
valid-r2:0.51007		
[250] train-rmse:8.67835	valid-rmse:8.29627	train-r2:0.54528
valid-r2:0.51073		
[260] train-rmse:8.63351	valid-rmse:8.29345	train-r2:0.54997
valid-r2:0.51106		
[270] train-rmse:8.59840	valid-rmse:8.28935	train-r2:0.55362
valid-r2:0.51155		
[280] train-rmse:8.57183	valid-rmse:8.28958	train-r2:0.55637
valid-r2:0.51152		
[290] train-rmse:8.54671	valid-rmse:8.28931	train-r2:0.55897
valid-r2:0.51155		
[300] train-rmse:8.51790	valid-rmse:8.28741	train-r2:0.56194

valid-r2:0.51177		
[310] train-rmse:8.48728	valid-rmse:8.28676	train-r2:0.56508
valid-r2:0.51185		
[320] train-rmse:8.45935	valid-rmse:8.28654	train-r2:0.56794
valid-r2:0.51188		
[330] train-rmse:8.43960	valid-rmse:8.28397	train-r2:0.56995
valid-r2:0.51218		
[340] train-rmse:8.41649	valid-rmse:8.28167	train-r2:0.57231
valid-r2:0.51245		
[350] train-rmse:8.39454	valid-rmse:8.27978	train-r2:0.57453
valid-r2:0.51267		
[360] train-rmse:8.37122	valid-rmse:8.28102	train-r2:0.57689
valid-r2:0.51253		
[370] train-rmse:8.34339	valid-rmse:8.27881	train-r2:0.57970
valid-r2:0.51279		
[380] train-rmse:8.31823	valid-rmse:8.27762	train-r2:0.58223
valid-r2:0.51293		
[390] train-rmse:8.29572	valid-rmse:8.27565	train-r2:0.58449
valid-r2:0.51316		
[400] train-rmse:8.27208	valid-rmse:8.27303	train-r2:0.58686
valid-r2:0.51347		
[410] train-rmse:8.24628	valid-rmse:8.26927	train-r2:0.58943
valid-r2:0.51391		
[420] train-rmse:8.22016	valid-rmse:8.26749	train-r2:0.59203
valid-r2:0.51412		
[430] train-rmse:8.18886	valid-rmse:8.26454	train-r2:0.59513
valid-r2:0.51447		
[440] train-rmse:8.16603	valid-rmse:8.26397	train-r2:0.59738
valid-r2:0.51453		
[450] train-rmse:8.13853	valid-rmse:8.26562	train-r2:0.60009
valid-r2:0.51434		
[460] train-rmse:8.11520	valid-rmse:8.26616	train-r2:0.60238
valid-r2:0.51428		
[470] train-rmse:8.09094	valid-rmse:8.26538	train-r2:0.60475
valid-r2:0.51437		
[480] train-rmse:8.06924	valid-rmse:8.26371	train-r2:0.60687
valid-r2:0.51456		
[490] train-rmse:8.04953	valid-rmse:8.26259	train-r2:0.60879
valid-r2:0.51470		
[500] train-rmse:8.02556	valid-rmse:8.26429	train-r2:0.61111
valid-r2:0.51450		
[510] train-rmse:8.00515	valid-rmse:8.26562	train-r2:0.61309
valid-r2:0.51434		
[520] train-rmse:7.98177	valid-rmse:8.26275	train-r2:0.61535
valid-r2:0.51468		
[530] train-rmse:7.96485	valid-rmse:8.26534	train-r2:0.61698
valid-r2:0.51437		
Stopping. Best iteration:		

```
[489]    train-rmse:8.05238    valid-rmse:8.26128    train-r2:0.60851  
valid-r2:0.51485
```

Predict test_df using XGBoost

```
[21]: p_test = clf.predict(d_test)
```

```
[22]: sub = pd.DataFrame()  
sub['ID'] = id_test  
sub['y'] = p_test  
sub.to_csv('test_df.csv', index = False)  
sub.head()
```

```
[22]:
```

	ID	y
0	1	82.865776
1	2	97.628395
2	3	83.197395
3	4	77.039124
4	5	112.527901

```
[ ]:
```